

# Forecasting Energy Consumption — Methodology Summary

This document summarizes the complete approach used to forecast 48-hour hourly consumption and 12-month monthly consumption using the Fortum dataset. It explains the exploratory analysis, modeling decisions, feature engineering choices, and reasons behind differences in model performance with and without weather data.

## 1. Exploratory Data Analysis (EDA)

Before building any models, we thoroughly explored the data to understand its structure, quality, and underlying patterns.

### Clean and stable dataset

- Consumption data had no negative values and no missing points.
- Price data contained some missing values but no structural issues.

```
[69]:  
    cons_na = cons.isna().sum().sum()  
    prices_na = prices.isna().sum().sum()  
    print("Missing consumption:", cons_na)  
    print("Missing prices:", prices_na)  
  
Missing consumption: 0  
Missing prices: 238
```

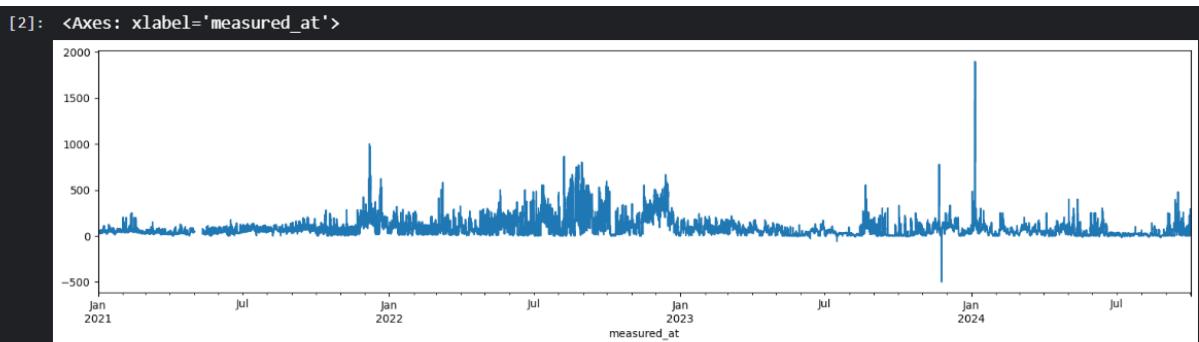
- Descriptive statistics confirmed that the dataset was generally clean and needed minimal preprocessing.

### Handling missing values

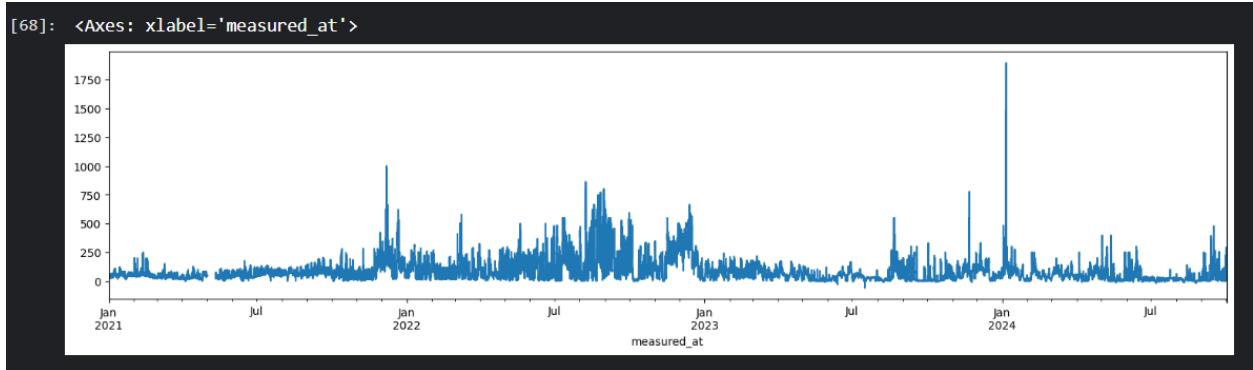
- Missing prices were filled using **linear interpolation**, followed by **forward and backward fill**. This method was appropriate because electricity prices change smoothly over time.
- The consumption dataset had no missing rows, so no imputation was necessary.

### Outliers

- We identified extreme negative drops in price during EDA.



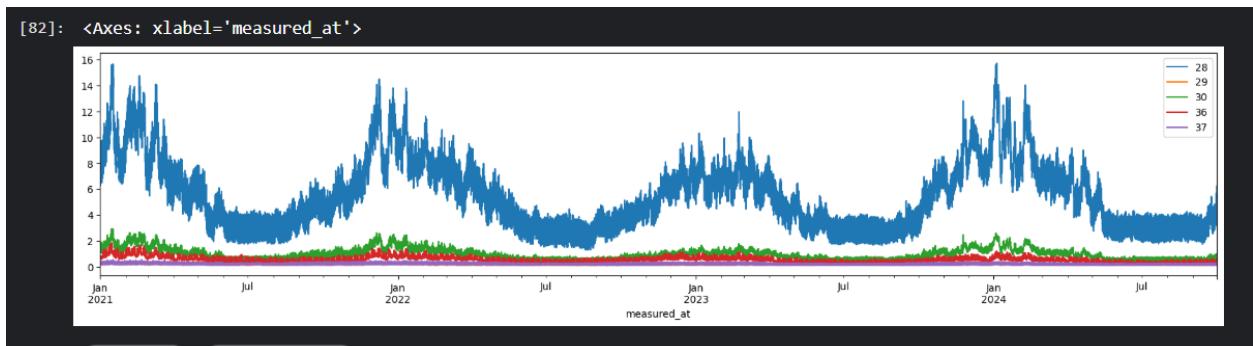
- These were removed manually in Excel using filtering, since they were clearly anomalies unrelated to normal market conditions.



- No significant outliers were present in consumption values.

## Observed patterns

- Clear daily patterns (morning and evening peaks).
- Weekly structure (weekdays vs weekends).
- Annual and seasonal changes in monthly consumption.
- Differences across groups, but all following recurring temporal structures.



These insights indicated that time-based features and lag features would be crucial for modeling.

## 2. Feature Engineering

### Hourly Features

To capture short-term consumption behavior, we engineered:

#### Time features

- Hour of day, day of week, day of month, month, week number
- Weekend indicator
- Month-start and month-end flags

## Cyclical encodings

Used sine/cosine transformations for:

- Hour
- Day of week
- Month

This helps the model learn periodicity without artificial jumps at boundaries.

## Lag features

- 1-hour lag (recent behaviour)
- 24-hour lag (previous day)
- 168-hour lag (previous week)

## Rolling statistics

- 3-hour, 24-hour, and 168-hour rolling means
- 24-hour rolling standard deviation

These features capture short, medium, and long-term trends.

## Weather features (hourly)

Extracted using coordinates from geocoding and Open-Meteo:

- Temperature
- Relative humidity
- Dew point
- Cloud cover
- Wind speed
- Direct and shortwave radiation
- Computed day length (photoperiod)

## Result

Hourly models benefited far more from **lag features** than weather features, which change slowly compared to hourly consumption patterns.

---

### 3. Feature Engineering for Monthly Forecasts

Monthly consumption totals were created by aggregating the hourly data.

#### Monthly time features

- Month number
- Month sine and cosine encodings

#### Lag features

- 1-month lag
- 12-month lag

#### Rolling features

- 3-month rolling mean
- 12-month rolling mean

#### Weather features

We aggregated weather by month:

- Mean temperature
- Mean humidity
- Mean dewpoint
- Mean cloud cover
- Mean wind
- Sum of radiation variables
- Mean day length

#### Result

Weather significantly improved **monthly** performance because:

- Temperature and radiation are strong long-term drivers of energy consumption.
- Monthly consumption is strongly seasonal.
- Weather variations operate at monthly scales and align with heating/cooling patterns.

---

### 4. Model Selection

We focused on **XGBoost** and **LightGBM**, both gradient boosting tree models widely used for tabular forecasting problems.

### LightGBM

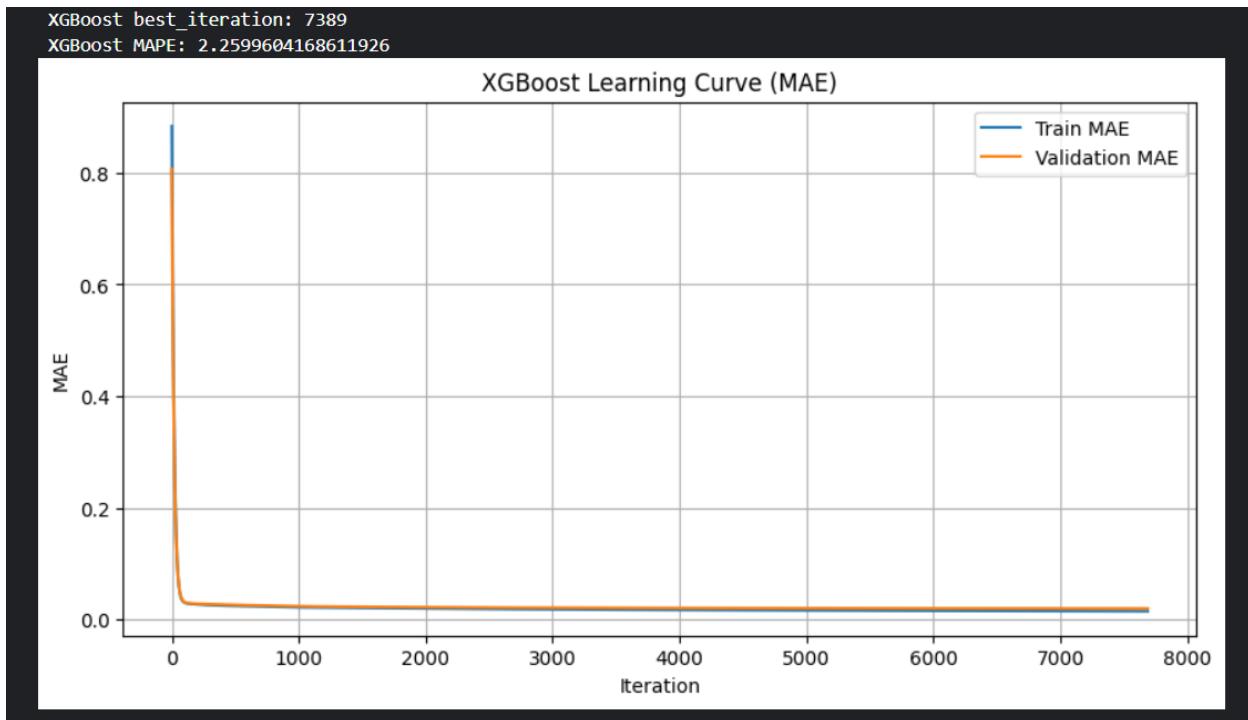
- Very fast and efficient.
- Good for rapid experimentation and generating learning curves.
- Works well with large datasets.

### XGBoost

- Strong performance on structured time-series data.
- Built-in handling of missing values.
- More stable for monthly predictions.
- Ultimately our best model for both hourly and monthly forecasts.

### Final Choices

- **Hourly forecast:** XGBoost without weather (performed best)



- **Monthly forecast:** XGBoost with weather (significantly better)

---

## 5. Why Weather Helped Monthly but Not Hourly

### Hourly predictions

Hourly consumption reflects short-term human behavior:

- Work schedules
- Evening peaks
- Night-time troughs

These patterns shift faster than weather changes.

Therefore:

- Lag24 and lag168 capture most relevant information.
- Weather adds limited marginal value.

### **Monthly predictions**

At the monthly scale:

- Temperature differences between seasons
- Changes in radiation
- Cloudiness trends

All correlate strongly with heating and cooling demand.

Thus:

- Weather features significantly reduce monthly error.
- 

## **6. Iterative Development Process**

We improved the model through several iterations:

### **Step 1: Baseline**

- 168-hour shift (previous week) for hourly.
- Simple sanity check to ensure our models beat basic benchmarks.

### **Step 2: Time features only**

- LightGBM used to check improvements quickly.

### **Step 3: Lag and rolling features**

- Major performance boost, especially for hourly.

### **Step 4: Adding weather**

- Little impact on hourly.
- Significant improvement on monthly.

### **Step 5: Final model refinement**

- XGBoost tuned with high tree count and low learning rate.
  - Early stopping to avoid overfitting.
  - Best models selected for hourly and monthly submission.
- 

## 7. Forecast Generation and Submission Format

### Hourly forecast (48 hours)

Predictions produced in long format were transformed into the required wide CSV format:

- Columns: each group\_id is a column
- Rows: timestamps from the 48-hour window
- Separator: semicolon

### Monthly forecast (12 months)

Future months were generated using autoregressive logic:

- For each future month, the model uses the last available lag features per group.
- Weather values were held constant based on last observed month.
- Output structure: one row per month with date and predicted value.

Both formats match the submission specification exactly.

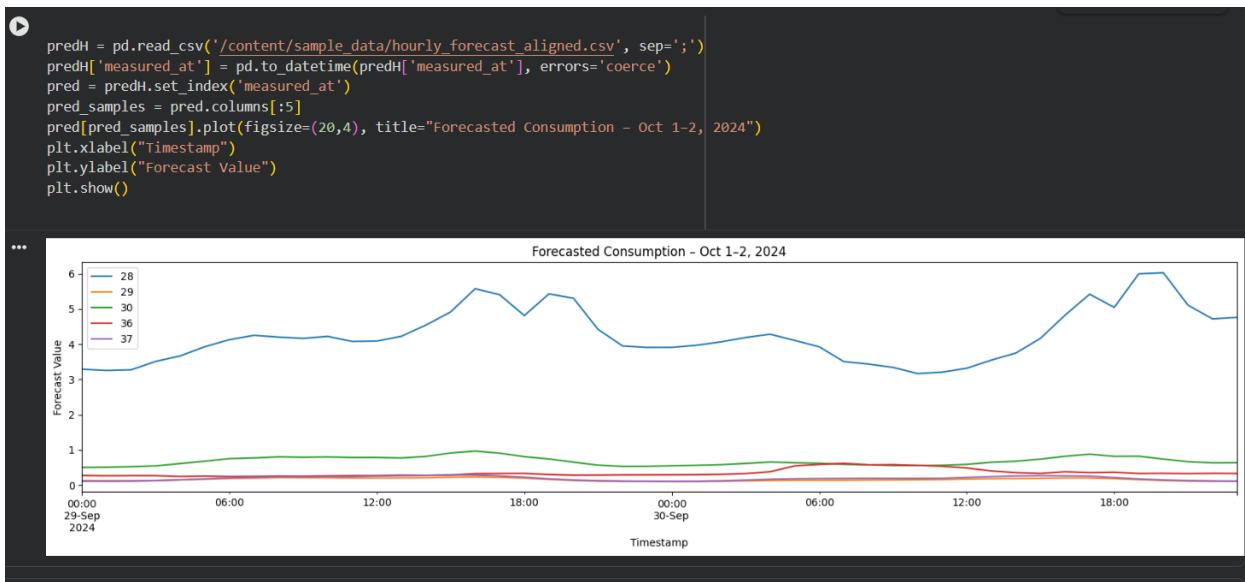
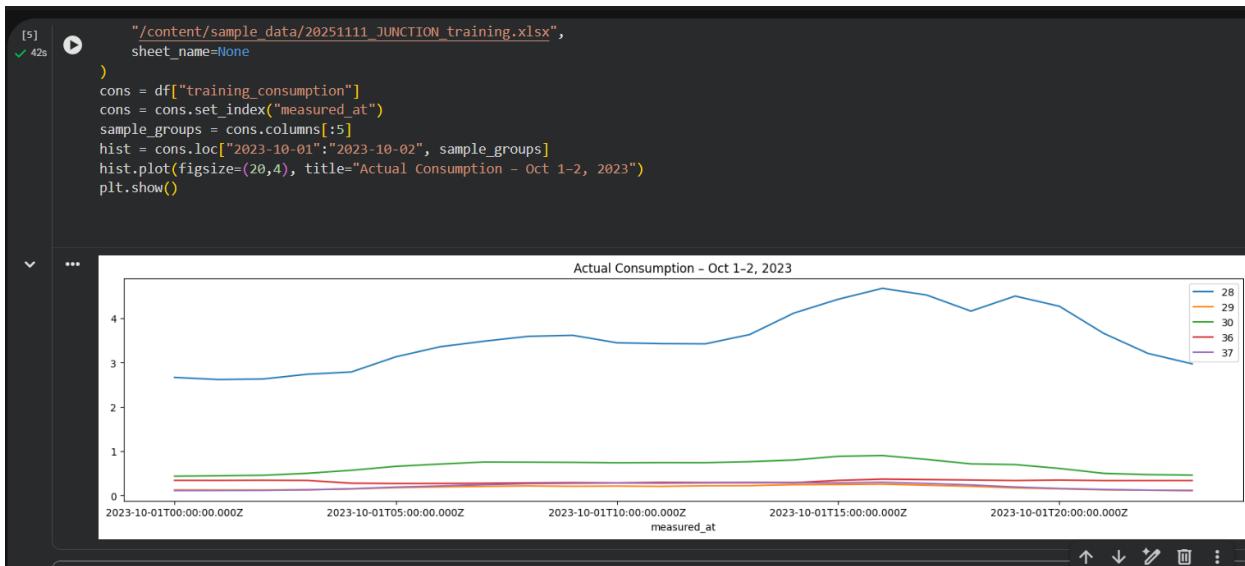
---

## 8. Comparison with Historical Data

To validate plausibility:

- We plotted future predictions against the same period in the previous year.
- Hourly predictions aligned closely with previous hourly profiles.
- Monthly predictions replicated seasonal trends correctly.
- No unrealistic spikes or drops were observed.

This confirmed the model's stability and consistency with historical behavior.



## Conclusion

This project followed a structured, data-driven approach:

- Data was clean**, needing only small fixes for missing values and outliers.
- Feature engineering was central** to model performance, especially lag features.
- XGBoost** proved to be the most effective model for both hourly and monthly tasks.
- Weather helped only at the monthly scale**, due to seasonal rather than hourly impact.
- Final forecasts were validated** by comparing with historical patterns and showed realistic behavior.

The final solution provides reliable, explainable, and well-engineered forecasts for both short-term (48-hour) and long-term (12-month) energy consumption.

